

Random Subset Methods for Forecast Combination: Theory, Monte Carlo Evidence, and Empirical Applications

Summary Report

Boris Kozyrev

June 2026

Contents

1	Theoretical Framework	3
1.1	Common-Loading Model	3
1.2	RSM Population Risk	3
1.3	Average-Subset Risk: First-Order Expansion and the Square-Root Rule . .	4
1.4	Infinite-Bagged Risk under Small Heterogeneity: the Cube-Root Rule . . .	4
1.5	General Heterogeneity and the Power-Law Approximation	5
2	Exogenous Monte Carlo (Known Σ_e)	6
2.1	Design	6
2.2	Main Results	7
2.3	T -Scaling and the Cube-Root Law	8
2.4	T -Scaling at $H = 6$: The Cube-Root Law Detected	8
3	Empirical Application: FRED-MD GDP Nowcasting	9
3.1	Data and Setup	9
3.2	Results: Bivariate RSM	9
3.3	Results: AR+X RSM	10
3.4	Connection to the Monte Carlo Evidence	11
4	SPF: The Near-Optimal Case	11
4.1	Setup and Motivation	11
4.2	Results: Hard to Beat the Mean	12

5 Overall Alignment: Theory, MC, and Empirics	13
A Data Sources	13
B m-Scaling: Pool Size and Optimal k^*	14
C Software and Replication	14

1 Theoretical Framework

1.1 Common-Loading Model

The baseline model for forecast error covariance is the *common-loading* (compound-symmetry) structure:

$$\Sigma_e = q \iota_m \iota_m' + D, \quad D = \text{diag}(\sigma_1^2, \dots, \sigma_m^2), \quad (1)$$

where ι_m is an m -vector of ones, $q \geq 0$ is a common variance component shared by all m forecasts, and σ_i^2 is the idiosyncratic variance of forecaster i . Denote precision weights $\tau_i = \sigma_i^{-2}$, $\bar{\tau} = m^{-1} \sum_i \tau_i$, and $A_S = \sum_{i \in S} \tau_i$ for a subset S . Write $\tau_i = \bar{\tau}(1 + \delta_i)$ with $\bar{\delta} = 0$ and heterogeneity measure $V_\delta = m^{-1} \sum_i \delta_i^2$.

The infeasible minimum-variance combination weight is $w_i^* = (\Sigma_e^{-1} \iota)_i / (\iota' \Sigma_e^{-1} \iota)$, with population risk $Q^{\text{opt}} = 1 / (\iota' \Sigma_e^{-1} \iota)$. For $q = 0$, $w_i^* = \tau_i / (m \bar{\tau})$ and $Q^{\text{opt}} = 1 / (m \bar{\tau})$.

1.2 RSM Population Risk

A Random Subset Method (RSM) draw: sample a size- k subset S uniformly at random from $\{1, \dots, m\}$, estimate sum-to-one constrained OLS weights $w_{i,S} = \tau_i / A_S$ on S , and form the subset forecast. Average G such draws to get $\hat{f}_{k,G}^{\text{bag}}$.

Average-subset risk.

$$Q_k^{\text{sub}} = q + \mathbb{E}_S \left[\frac{1}{A_S} \right]. \quad (2)$$

Infinite-bagged risk. Let $\bar{v}_{i,k} = \mathbb{E}_S[w_{i,S} \mathbf{1}(i \in S)]$ be the average weight placed on forecaster i . Then

$$Q_k^{\text{bag},\infty} = \bar{v}_k' \Sigma_e \bar{v}_k. \quad (3)$$

Weight-deviation representation. The excess risk of the infinite-bagged estimator over the infeasible optimum is a quadratic form in the deviation of the average RSM weight from the optimal weight:

$$Q_k^{\text{bag},\infty} - Q^{\text{opt}} = (\bar{v}_k - w^*)' \Sigma_e (\bar{v}_k - w^*) \geq 0. \quad (4)$$

This representation is key: *the functional form of the excess risk is determined entirely by how the average RSM weight $\bar{v}_k$ deviates from w^** . The deviation vanishes as $k \rightarrow m$ (full pool), and its dependence on k drives the optimal subset size.

Finite-bagging bridge. For G draws the expected loss is

$$Q_k^{\text{bag},G} = Q_k^{\text{bag},\infty} + \frac{1}{G} \left(Q_k^{\text{sub}} - Q_k^{\text{bag},\infty} \right). \quad (5)$$

This identity is exact (no approximation), and was verified in simulation to machine precision (max error $< 9 \times 10^{-15}$). In particular, $Q_k^{\text{bag},\infty} \leq Q_k^{\text{sub}}$ for all k : infinite bagging is weakly better than a single subset. The gain from bagging grows with G .

1.3 Average-Subset Risk: First-Order Expansion and the Square-Root Rule

The average-subset estimator uses a single random subset; no averaging takes place. Its excess risk comes from $\mathbb{E}_S[1/A_S]$ in (2). Under homogeneous precision ($\tau_i = \bar{\tau}$), $A_S = k\bar{\tau}$ exactly, so $\mathbb{E}_S[1/A_S] = 1/(k\bar{\tau})$ and

$$Q_k^{\text{sub}} - Q^{\text{opt}} \approx \frac{1}{k\bar{\tau}} - \frac{1}{m\bar{\tau}} = \frac{z_k}{\bar{\tau}}, \quad z_k \equiv \frac{1}{k} - \frac{1}{m}. \quad (6)$$

The excess risk is *linear* in z_k (effectively $\alpha = 1$).

Including finite-sample weight-estimation inflation $(T-1)/(T-k)$, the expected in-sample loss is

$$L_k^{\text{sub}} \approx \left(Q^{\text{opt}} + \frac{z_k}{\bar{\tau}} \right) \frac{T-1}{T-k}. \quad (7)$$

Setting $\partial L_k^{\text{sub}} / \partial k = 0$ and approximating $z_k \approx 1/k$ for $k \ll m$ yields

$$k_{\text{sub}}^* \approx \left(\frac{(T-1)}{\bar{\tau} Q^{\text{opt}}} \right)^{1/2} \propto T^{1/2}. \quad (8)$$

The average-subset estimator has a *square-root* optimal subset size.

1.4 Infinite-Bagged Risk under Small Heterogeneity: the Cube-Root Rule

Bagging replaces the random single-subset weight $w_{i,S}$ with its expectation $\bar{v}_{i,k}$, concentrating the weight and reducing variance. To find the functional form of the excess risk (4), we compute the weight deviation $\Delta_{i,k} = \bar{v}_{i,k} - w_i^*$.

Weight deviation under small heterogeneity. With $\tau_i = \bar{\tau}(1 + \delta_i)$ and $V_\delta \ll 1$, a first-order Taylor expansion of τ_i/A_S around the homogeneous case gives

$$\Delta_{i,k} = \bar{v}_{i,k} - w_i^* \approx -\frac{\delta_i z_k}{m\bar{\tau}}. \quad (9)$$

The deviation is proportional to z_k : it vanishes as $k \rightarrow m$ (the bagged estimator recovers the optimum when all forecasters are included) and is scaled by the idiosyncratic heterogeneity δ_i .

Excess risk as a quadratic in z_k . Substituting (9) into (4) and using $\Sigma_e = q\iota' + D$:

$$Q_k^{\text{bag},\infty} - Q^{\text{opt}} \approx \Delta' D \Delta = \frac{z_k^2}{m^2 \bar{\tau}^2} \sum_i \sigma_i^2 \delta_i^2 = C_\delta \left(\frac{1}{k} - \frac{1}{m} \right)^2. \quad (10)$$

The common-loading term $q(\iota' \Delta)^2$ vanishes exactly because $\sum_i \delta_i = 0$. The coefficient $C_\delta = V_\delta / (m \bar{\tau}^3)$ captures the dispersion of precision weights.

The excess risk is *quadratic* in z_k (i.e. $\alpha = 2$ in a log-log sense), in contrast to the linear $\alpha = 1$ decay for Q^{sub} . Bagging effectively *squares* the decay rate of excess risk with k .

Optimal k for the bagged estimator. The total loss of the infinite-bagged estimator, including finite-sample inflation, is

$$L_k^{\text{bag}} \approx \left[Q^{\text{opt}} + C_\delta \left(\frac{1}{k} - \frac{1}{m} \right)^2 \right] \frac{T-1}{T-k}. \quad (11)$$

Setting $\partial L_k^{\text{bag}} / \partial k = 0$ and approximating $z_k \approx 1/k$ for $k \ll m$ yields the *cube-root optimal subset size*:

$$k_{\text{bag}}^* \approx \left(\frac{2C_\delta T}{Q^{\text{opt}}} \right)^{1/3} \propto T^{1/3}. \quad (12)$$

The T -scaling exponent drops from $\frac{1}{2}$ (single subset) to $\frac{1}{3}$ (infinite bag): bagging makes larger subsets optimal relative to the average-subset rule.

1.5 General Heterogeneity and the Power-Law Approximation

When heterogeneity is large, equation (10) no longer holds and the weight deviations $\Delta_{i,k}$ do not scale as $\delta_i z_k$. Empirically, the log-log regression

$$\log(Q_k^{\text{bag},\infty} - Q^{\text{opt}}) \approx \hat{\alpha} \log z_k + \text{const} \quad (13)$$

yields an *effective log-log slope* $\hat{\alpha}$ that lies in $[1, 2]$ depending on the degree of heterogeneity. Small heterogeneity gives $\hat{\alpha} \approx 2$ (the small-het formula, eq. (10)); high heterogeneity produces $\hat{\alpha}$ closer to 1.

Remark 1 (Effective vs. exact exponent). $\hat{\alpha}$ is an OLS slope of (13) over the chosen k -grid, not an exact mathematical exponent. In the dominant-forecaster limit ($\tau_1 \gg \tau_j$ for $j \neq 1$), the exact curve is $R_k = C'(1 - k/m)^2$, which is quadratic in $(1 - k/m)$ but

not a pure power of $z_k = (m - k)/(km)$. Near $k = m$, it behaves locally as $C'm^2 z_k^2$, but an OLS regression over the full grid gives a slope near 1 because $(1 - k/m)^2$ and z_k have different curvatures over the empirically relevant range of k . We therefore call $\hat{\alpha}$ the effective slope; the bounds $\hat{\alpha} \in [1, 2]$ describe the OLS projection, not pointwise power-law identities. Simulations consistently recover $\hat{\alpha} \in [1, 2]$ across all heterogeneity levels studied, with $\hat{\alpha}$ increasing from ≈ 1 at low H to ≈ 1.7 at high H (Table 1).¹

Minimising the generalised loss $L_k \approx (A + Cz_k^{\hat{\alpha}})(T - 1)/(T - k)$ using the effective slope $\hat{\alpha}$ gives

$$k^* \approx \left(\frac{C\hat{\alpha}T}{A} \right)^{1/(\hat{\alpha}+1)}, \quad (14)$$

so the effective T -scaling exponent $1/(\hat{\alpha} + 1)$ ranges from $\frac{1}{3}$ (small-het, $\hat{\alpha} = 2$) to $\frac{1}{2}$ (strongly heterogeneous, $\hat{\alpha} = 1$). The cube-root rule (12) is the *lower bound* on the effective scaling exponent.

Remark 2 (Effective scaling caveat). *Formula (14) uses the fitted power-law approximation and should be read as an effective scaling rule. It accurately describes the optimal subset size when the log-log fit is tight (as confirmed empirically for all four DGPs in Table 1 with $R^2 > 0.97$). Under the exact dominant-forecaster curve the scaling exponent is approximately $1/(\hat{\alpha} + 1)$ with the effective $\hat{\alpha}$, giving the same practical guidance.*

In both cases k^* is increasing in T and in $\hat{\alpha}$ (higher heterogeneity requires a larger subset to cover the precision-weight distribution).

2 Exogenous Monte Carlo (Known Σ_e)

2.1 Design

The exogenous experiment takes Σ_e as given (no estimation). The DGP follows the common-loading structure with $m = 50$, $T = 240$ (baseline). Forecast errors are generated as

$$x_{it} = B_{\text{scale}} \cdot F_t + H_{\text{scale}} \cdot \varepsilon_{it}, \quad (15)$$

where $F_t \sim N(0, 1)$ is a common factor and ε_{it} are idiosyncratic errors. The parameter $B \in \{3, 8\}$ controls the common variance q and $H \in \{1.05, 6\}$ controls the spread of individual precisions τ_i (heterogeneity). Simulations use $N_{\text{rep}} = 500$ replications and $G = 800$ bags. Key theory quantities (Q^{opt} , C_δ , k^* formulae) are computed analytically from the true Σ_e .

¹The dominant-forecaster weight deviation for $j \neq 1$ is $\bar{v}_{j,k} = P(j \in S) \cdot P(1 \notin S \mid j \in S) \cdot (1/k) = (k/m) \cdot (m - k)/(m - 1) \cdot (1/k) = (m - k)/(m(m - 1))$, so $\Delta_{j,k} = \bar{v}_{j,k} - w_j^* \approx (m - k)/(m(m - 1))$ since $w_j^* \rightarrow 0$ as $\tau_1 \rightarrow \infty$. The resulting $R_k \approx (m - 1)\Delta_{j,k}^2/\tau_j = (m - k)^2/(m^2(m - 1)\tau_j)$ is *not* a power of z_k globally; locally near $k = m$ it is $\propto z_k^2$.

2.2 Main Results

Table 1 summarises the four DGP combinations, using all available population-level and finite-sample statistics. The column “ $k^*(\text{grid})$ ” is the minimiser of $Q_k^{\text{bag},\infty}$ over $k \in \{1, \dots, m\}$; “ $k^*(\text{emp})$ ” is the minimiser of the finite-sample MSPE over 500 replications; “Gain” is the percentage improvement of the optimal RSM over the equal-weighted mean.

Table 1: Exogenous MC: $m = 50$, $T = 240$, 500 replications, 800 bags.

B	H	$k^*(\text{grid})$	$k^*(\text{emp})$	Gain vs EW	$\hat{\alpha}$	R^2	Approx. err
3	1.05	1	3	0.1%	1.12	0.98	90%
8	1.05	1	3	0.1%	1.10	0.98	90%
3	6	3	6	12.7%	1.72	0.98	75%
8	3.5*	4	8	14.9%	1.70	0.99	77%

* Target $H = 6$; parameterisation achieved $H_{\text{actual}} = 3.51$.

$\hat{\alpha}$: OLS slope in $\log(Q_k^{\text{bag},\infty} - Q^{\text{opt}}) \sim \hat{\alpha} \log z_k$; Approx. err: mean $|R_k - C_\delta z_k^2|/R_k$ over k .

Theory alignment. Several predictions are confirmed:

- **Finite-bagging bridge:** equation (5) holds to machine precision ($< 9 \times 10^{-15}$) for all k and all DGPs.
- **Cube-root formula:** $k^*(\text{formula}) = k^*(\text{grid})$ in all four DGPs, confirming that (12) locates the correct minimiser even when the small-het approximation fails by a factor of 5–100 $\times$.
- **Power-law:** $\hat{\alpha} \in [1, 2]$ with $R^2 > 0.97$ in all cases. $\hat{\alpha}$ increases with H (from ≈ 1.1 at $H \approx 1$ to ≈ 1.7 at $H \approx 6$), consistent with the theoretical bounds. It depends on H (heterogeneity) but *not* on B (common factor).
- **Gains grow with heterogeneity:** 0.1% at $H = 1.05$ vs 12–15% at $H = 3.5$ –6. RSM is most valuable when forecasters differ substantially in quality.

Small-het approximation failure. The approximation error (90% at $H = 1.05$, 75–77% at $H = 3.5$ –6) is large. The true excess risk exceeds $C_\delta z_k^2$ substantially at moderate k , as the Taylor expansion of $\mathbb{E}_S[1/A_S^2]$ involves higher-order heterogeneity moments. Despite this, the cube-root formula gives the exact minimiser because the leading bias term in the approximation is constant in k and does not affect the location of the minimum.

2.3 T -Scaling and the Cube-Root Law

The T -scaling experiment varies $T \in \{80, 240, 640, 1000\}$ at $B = 3$, $H = 1.05$. Table 2 shows the continuous cube-root formula value, the grid-search optimal, and the empirical optimum.

Table 2: Exogenous T -scaling: $B = 3$, $H \approx 1.05$, $m = 50$.

T	$k^*(\text{formula, cont.})$	$k^*(\text{grid, bag}^\infty)$	$k^*(\text{emp})$	Gain
80	0.57	1	2	0.02%
240	0.82	1	2	0.02%
640	1.13	1	4	0.06%
1000	1.31	1	3	0.03%

At $H = 1.05$, the population excess risk is virtually zero for any k (gains $< 0.1\%$), so the loss function is essentially flat and the population optimum is $k^* = 1$ throughout. The cube-root formula correctly predicts this: the continuous value ranges from 0.57 to 1.31, confirming that $k^* = 1$ is optimal at every T studied. The cube-root growth rate is present but invisible because the signal (heterogeneity) is negligible.

The empirical $k^* = 2\text{--}4$ reflects simulation noise on a near-flat loss surface, not a structural deviation from the theory. Detecting the cube-root law requires sufficient heterogeneity ($H \gg 1$); the next subsection confirms this at $H = 6$.

2.4 T -Scaling at $H = 6$: The Cube-Root Law Detected

To confirm that the cube-root scaling law is present but masked at low heterogeneity, we repeat the T -scaling experiment at $H = 6$ ($B = 3$, $m = 50$, 500 replications). Table 3 reports the results.

Table 3: Exogenous T -scaling: $B = 3$, $H \approx 6.08$, $m = 50$, 500 reps.

T	$k^*(\text{formula, cont.})$	$k^*(\text{grid, bag}^\infty)$	$k^*(\text{emp})$	Gain
80	2.0	2	4	12.4%
240	2.9	3	6	12.1%
640	4.0	4	8	12.8%
1000	4.6	5	10	12.8%

The log-log regression of $k^*(\text{grid})$ on T gives slope $\hat{\beta} = 0.352$ ($R^2 = 0.994$) — essentially exactly $\frac{1}{3}$. The result for $k^*(\text{emp})$ is identical ($\hat{\beta} = 0.352$, $R^2 = 0.994$). In contrast to the $H = 1.05$ case (where gains are $< 0.1\%$ and $k^* = 1$ throughout), at $H = 6$ gains are approximately constant at 12–13% across all T , and the cube-root scaling is clearly detectable. This confirms that the law is structural, not an artefact of heterogeneity level: it holds whenever heterogeneity is non-trivial.

Connection to the formula. The continuous cube-root formula (12) predicts $k^* \in \{2.0, 2.9, 4.0, 4.6\}$ for $T \in \{80, 240, 640, 1000\}$, matching the grid search exactly after rounding. The effective $\hat{\alpha}$ at $H = 6$ is 1.72 (Table 1), giving a predicted scaling exponent of $1/(1.72 + 1) = 0.368$, consistent with the empirical 0.352.

3 Empirical Application: FRED-MD GDP Nowcasting

3.1 Data and Setup

We use 102 real-time quarterly vintages of FRED-MD (September 1999 – March 2025), excluding Q4 2003 (missing data). At each vintage, the estimation sample covers approximately 1970Q1 to the forecast origin, giving $T_{\text{eff}} \approx 79\text{--}183$ quarters (median ≈ 100). The predictor pool contains $m \approx 126$ monthly FRED-MD series, aggregated to quarterly frequency.

Individual forecasts (bivariate). For each predictor series j and time t , we estimate

$$y_t = \alpha_j + \beta_j x_{j,t} + \varepsilon_t, \quad (16)$$

where y_t is log-differenced real GDP growth. This yields m individual forecasts per period.

Individual forecasts (AR+X). Following Elliott et al. (2013), we also produce forecasts from

$$y_t = \alpha_j + \varphi_1 y_{t-1} + \cdots + \varphi_p y_{t-p} + \beta_j x_{j,t} + \varepsilon_t, \quad (17)$$

where p is chosen by AIC at each t .

RSM combination. At each vintage and period t , RSM samples k individual forecasts and runs sum-to-one constrained OLS. We average $B = 1000$ draws. For robustness we sweep $k \in \{2, \dots, 55\}$ with $B = 200$ draws.

Benchmarks. Equal-weighted mean, Lasso, Ridge, Random Forest.

3.2 Results: Bivariate RSM

RMSFE by k . The $\text{RMSFE}(k)$ curve is U-shaped. Key values:

- Full sample: $k^*(\text{RMSFE}) = 19$, $\min \text{RMSFE} = 0.662$. At $k = \sqrt{T} \approx 9$: $\text{RMSFE} = 0.762$.
- No-COVID: $k^*(\text{RMSFE}) = 21$, $\min \text{RMSFE} = 0.499$. At $k = 9$: $\text{RMSFE} = 0.529$.

Table 4: FRED-MD GDP: bivariate RSM, full sample (102 quarters).

Model	MAFE	Sig.	RMSFE	Sig.
Mean	0.54		1.10	
RSM	0.45	**	0.76	*
Lasso	0.51		0.73	*
Ridge	0.49	*	0.98	
Random Forest	0.55		1.24	
$k = \lfloor \sqrt{T} \rfloor \approx 9$. * $p < 0.10$, ** $p < 0.05$ (DM test vs. Mean).				

Table 5: FRED-MD GDP: bivariate RSM, excluding COVID-19 (Q2–Q3 2020).

Model	MAFE	Sig.	RMSFE	Sig.
Mean	0.42		0.60	
RSM	0.38	**	0.53	**
Lasso	0.45		0.58	
Ridge	0.39	**	0.54	**
Random Forest	0.41		0.57	*

3.3 Results: AR+X RSM

Table 6: FRED-MD GDP: AR+X RSM, full sample.

Model	MAFE	Sig.	RMSFE	Sig.
Mean	0.56		1.27	
RSM (AR+X)	0.45	**	0.76	*
Lasso	0.52		0.77	
Ridge	0.50	**	1.02	*
Random Forest	0.57		1.40	

Bivariate vs. AR+X comparison.

Sample	Bivariate		AR+X	
	k^*	min RMSFE	k^*	min RMSFE
Full	19	0.662	19	0.617
No-COVID	21	0.499	16	0.506

Adding AR lags lowers the full-sample minimum RMSFE by 6.8%. In the no-COVID sample, the benefit disappears (AR lags add estimation noise without large cyclical episodes to track). The AR+X $k^* = 16$ (no-COVID) is closer to the MC range of 11–14 than the bivariate $k^* = 21$.

Table 7: FRED-MD GDP: AR+X RSM, excluding COVID-19.

Model	MAFE	Sig.	RMSFE	Sig.
Mean	0.41		0.57	
RSM (AR+X)	0.39	*	0.52	*
Lasso	0.47		0.63	
Ridge	0.39	**	0.54	**
Random Forest	0.40		0.57	

3.4 Connection to the Monte Carlo Evidence

Table 8: Comparison of $k^*(\text{RMSFE})$ across MC and empirical settings.

Setting	T	$k^*(\text{RMSFE})$
GDP bivariate, full sample	≈ 100	19
GDP bivariate, no-COVID	≈ 100	21
GDP AR+X, full sample	≈ 100	19
GDP AR+X, no-COVID	≈ 100	16

The empirical k^* (16–21 at $T \approx 100$) is larger than what the exogenous population theory predicts ($k^* \approx 1$ –5 for $T = 240$), for two structural reasons:

1. **Underpowered covariance estimation:** GDP has $m \approx 126$ predictors with $T \approx 100$ observations ($m/T \approx 1.3$). Estimating combination weights is much harder than in the exogenous MC (where Σ_e is known). The additional estimation noise pushes k^* substantially rightward.
2. **Individual forecast noise:** each predictor’s individual forecast carries its own OLS estimation error (from the bivariate/AR+X regression). Lower signal-to-noise per individual forecast makes larger subsets more attractive.

The AR+X specification reduces but does not eliminate this gap: it improves individual forecast quality, which brings the no-COVID k^* from 21 to 16. Both remain well above the exogenous population optimum, which is characteristic of high-dimensional, low- T empirical settings.

4 SPF: The Near-Optimal Case

4.1 Setup and Motivation

The theory predicts that RSM gains are negligible when forecasters are near-homogeneous: at $H \approx 1$, the cube-root formula gives $k^* \approx 1$ and gains are $< 0.1\%$ (Table 1). The

Survey of Professional Forecasters (SPF) provides a direct test of this prediction: SPF respondents are trained professional economists using similar information sets, making near-homogeneity ($H \approx 1$) the natural prior.

We apply RSM directly to SPF responses, replicating the setup of Conflitti et al. (2015). The sample covers Q1 1985 – Q4 2024, with $N_t \approx 30$ –53 respondents per round (median 36). Each forecaster’s own submission is the individual forecast (no regression estimation). $B = 1000$ draws per round. Coverage constraint: $k \leq k_{\max}$ where $k_{\max}$ is the largest k at which $\geq 80\%$ of rounds have $N_t \geq k$.

Four target variables: real GDP growth ($h = 1$ and $h = 4$ quarters ahead) and CPI inflation ($h = 1$ and $h = 4$ quarters ahead).

4.2 Results: Hard to Beat the Mean

Table 9: SPF RSM vs. equal-weighted mean; OOS 1985Q1–2024Q4.

Series	h	Median N	k^*	$k = \sqrt{N}$	RMSFE (Mean)	RMSFE (RSM)
RGDP	1	36	3	6	6.405	6.392 (−0.2%)
RGDP	4	36	13	6	2.208	2.205 (−0.1%)
CPI	1	35	10	6	0.974	0.969 (−0.5%)
CPI	4	35	3	6	1.471	1.470 (−0.1%)

No DM tests significant at 10% level.

Gains are uniformly below 0.5% and insignificant. This confirms the main theoretical prediction: *it is hard to beat the equal-weighted mean when forecasters are near-homogeneous*. Several features of the SPF environment explain the small gains:

1. **Near-homogeneous precision.** Professional forecasters use similar models and information sets; heterogeneity $H \approx 1$. The MC shows that at $H = 1.05$ gains are $< 0.1\%$, matching the SPF results.
2. **No regression estimation noise.** SPF individual forecasts are direct submissions, without additional estimation error. In the FRED-MD setting, each individual forecast carries OLS estimation noise that RSM partially averages out. Removing that noise removes a channel through which RSM gains arise.
3. **Small pool.** At $N \approx 36$ respondents, distinct-subset diversity is limited. The combinatorial benefit of RSM grows with m ; at $m = 36$ it is much smaller than at $m = 126$ (FRED-MD) or $m = 50$ (MC).

Despite the small magnitude, the SPF results confirm the *direction* of all MC predictions: RSM never underperforms the mean; k^* is interior (not at the boundary); and the $\sqrt{N}$ rule of thumb overestimates k^* in the low-heterogeneity cases (RGDP $h = 1$, CPI $h = 4$) and underestimates it in the moderate-heterogeneity case (CPI $h = 1$).

5 Overall Alignment: Theory, MC, and Empirics

Table 10 provides a unified scorecard of the theoretical predictions across the three empirical settings.

Table 10: Theory–MC–Empirics alignment.

Prediction	Exog. MC	Empirical (GDP)	SPF
Finite-bagging bridge (exact)	✓	n/a	n/a
Cube-root formula = k^* (grid)	✓	×	n/a
Power-law $\hat{\alpha} \in [1, 2]$	✓	n/a	n/a
$\hat{\alpha}$ depends on H , not B	✓	n/a	n/a
k^* scales as $T^{1/3}$ at $H = 6$	✓	n/a	n/a
k^* scales as $T^{1/(\hat{\alpha}+1)}$	✓	~	n/a
RSM beats equal weights	✓	✓	✓
Near-homogeneous: gains < 1%	✓	n/a	✓
$\sqrt{T}$ rule of thumb for k^*	~	~	~
Gains grow with heterogeneity	✓	✓	✓
U-shaped RMSFE(k) in practice	n/a	✓	n/a

✓: confirmed; ~: partial; ×: not confirmed; n/a: not applicable.

Summary. The exogenous MC validates the core theoretical machinery: the exact finite-bagging bridge, the cube-root formula, and the power-law decay of excess risk all hold precisely. The key empirical finding is that heterogeneity (H) is the decisive parameter: at $H \approx 1$ (SPF, near-homogeneous forecasters) RSM gains are negligible and it is hard to beat the simple mean. At $H \gg 1$ (FRED-MD GDP, heterogeneous predictors) RSM produces statistically significant gains of 10–30%.

The FRED-MD k^* (16–21) exceeds the exogenous population prediction because estimation of combination weights is expensive when $m/T > 1$: the optimal k must balance approximation bias against the additional noise from fitting constrained OLS on each subset. The AR+X specification brings k^* closer to the MC range by improving individual forecast quality, confirming the structural interpretation.

A Data Sources

- **FRED-MD:** Real-time monthly vintages, September 1999 – March 2025. 103 vintages; Q4 2003 excluded (missing data). McCracken and Ng (2016).
- **FRED-QD:** Quarterly vintages matched to monthly by vintage date.
- **SPF:** Philadelphia Fed Survey of Professional Forecasters. RGDP file (columns RGDP1–RGDP5) and CPI file (columns CPI1–CPI5). Actuals: GDPC1 and CPI-AUCSL from FRED (downloaded June 2026).

B m -Scaling: Pool Size and Optimal k^*

The cube-root formula (12) scales with T but does not explicitly depend on m . In practice, larger pools require larger subsets because (i) the precision weight distribution is more spread out and (ii) estimating combination weights is noisier when m is large relative to T . To quantify the m -dependence, we run the endogenous MC at $T = 200$ fixed and $m \in \{20, 30, 50, 80, 120\}$ with paper-level settings ($N_{\text{rep}} = 500$, $K_{\text{RS}} = 500$, OOS length = 50, B4 coefficients).

Table 11: m -scaling: $k^*(\text{RMSFE})$ by pool size, $T = 200$ fixed.

Σ_e	$m = 20$	$m = 30$	$m = 50$	$m = 80$	$m = 120$	log-log slope
Toeplitz	8	8	12	17	22	0.611 ($R^2 = 0.96$, $p = 0.004$)
Factor	8	10	14	20	18*	0.512 ($R^2 = 0.91$, $p = 0.012$)

Factor $m = 120$: k^ dips to 18 from 20 at $m = 80$, likely grid-step noise (step size 4).

Log-log slope: OLS of $\log k^*$ on $\log m$. Gains stabilise at 10–11% (Toeplitz) and 4–7% (Factor) across m .

k^* rises sub-linearly with m for both Σ_e types, with log-log slopes of 0.61 (Toeplitz) and 0.51 (Factor) — well below 1 (linear) and above 0.33 (cube-root). This confirms that larger pools require larger subsets, with the rate between the $T^{1/3}$ and $T^{1/2}$ predictions. The Kan-Zhou correction factor $h(m, T) = (T-2)/(T-m-2)$ (which captures estimation noise from estimating Σ_e from m series) rises from 1.11 at $m = 20$ to 2.54 at $m = 120$, quantitatively consistent with the observed k^* trajectory for Toeplitz.

C Software and Replication

All code is in R (≥ 4.2). Key packages: `dplyr`, `glmnet`, `ranger`, `forecast`, `readxl`. All random number streams are seeded. Scripts:

- Exogenous MC: `rsm_common_loading_grid_weights...final.R`
- GDP bivariate: `GDP_constrained/GDP_all_k_results.R`
- GDP AR+X: `GDP_constrained/GDP_ar_bivariate_allk.R`
- SPF: `GDP_constrained/SPF_RSM_analysis.R`
- m -scaling MC: `rsm_m_scaling.R`

References

Conflitti, C., De Mol, C., & Giannone, D. (2015). Optimal combination of survey forecasts. *International Journal of Forecasting*, 31(4), 1096–1103.

- Elliott, G., Gargano, A., & Timmermann, A. (2013). Complete subset regressions. *Journal of Econometrics*, 177(2), 357–373.
- Kan, R., & Zhou, G. (2007). Optimal portfolio choice with parameter uncertainty. *Journal of Financial and Quantitative Analysis*, 42(3), 621–656.